

SCAFFOLD STRUCTURE

intent_extractor.py

Calls Ollama with a structured JSON schema to extract {goal, constraints, context, success_criteria} from every turn, then merges it with the previous turn's intent. So if the user says "make it faster" on turn 4, the system still knows what "it" refers to from turn 1.

outcome_feedback.py

Watched the next user message to infer whether the previous response worked. Rephrasing the same question = wrong. "Building on" language = right. Explicit frustration markers = wrong. No more length-heuristic nonsense.

session_tracker.py

Cross-turn memory the middleware can't see. Detects topic drift via Jaccard similarity on content words, repeat asks via string similarity, and frustration arcs by scoring escalating markers across the rolling window. Produces a recommendation that can override the middleware's route.

confidence_floor.py

The only module that can stop the agent from producing anything. Weighted score across four components; below the floor it refuses with an honest message rather than hallucinating intent. The thresholds are at the top and worth tuning once you have eval data.

bias_cooccurrence.py

8 named interaction effects with specific routing implications. Dunning-Kruger + confirmation bias together (double_certainty) is the worst case — compliance makes it worse, direct challenge gets rejected, so the route becomes a careful surface-the-gap challenge.

eval_harness.py

21 baseline cases, per-bias accuracy breakdown, latency measurement, classifier comparison mode. The most important thing here: add your own cases from real sessions. The baseline will have gaps specific to how you actually talk.